

class10.Rmd

Sara Herrera (PID:A59011948)

10/29/2021

1. Importing candy data

```
candy_file <- "candy-data.csv"
```

```
candy = read.csv(candy_file, row.names=1)  
head(candy)
```

```
##           chocolate fruity caramel peanutyalmondy nougat crispedricewafer  
## 100 Grand           1      0          1              0      0              1  
## 3 Musketeers        1      0          0              0      1              0  
## One dime            0      0          0              0      0              0  
## One quarter        0      0          0              0      0              0  
## Air Heads           0      1          0              0      0              0  
## Almond Joy          1      0          0              1      0              0  
##           hard bar pluribus sugarpercent pricepercent winpercent  
## 100 Grand          0  1          0          0.732      0.860  66.97173  
## 3 Musketeers        0  1          0          0.604      0.511  67.60294  
## One dime            0  0          0          0.011      0.116  32.26109  
## One quarter        0  0          0          0.011      0.511  46.11650  
## Air Heads           0  0          0          0.906      0.511  52.34146  
## Almond Joy          0  1          0          0.465      0.767  50.34755
```

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
## [1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
## [1] 38
```

2. What is your favourite candy?

Q3 & Q4. What is your favorite candy in the dataset and what is its winpercent value?

```
candy["Kit Kat", ]$winpercent
```

```
## [1] 76.7686
```

Q5. What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
## [1] 49.6535
```

```
library("skimr")  
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency: numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?


There's some variables/columns on a different scale, in percentages.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

Zero is NO and 1 is YES

Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent)
```



Q9. Is the distribution of winpercent values symmetrical?

No

Q10. Is the center of the distribution above or below 50%?

```
summary(candy$winpercent)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	22.45	39.14	47.83	50.32	59.86	84.18

Below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Higher ranked

```
chocolate <- candy[ as.logical(candy$chocolate), ]$winpercent
mean(chocolate)
```

```
## [1] 60.92153
```

```
fruity <- candy[ as.logical(candy$fruity), ]$winpercent
mean(fruity)
```

```
## [1] 44.11974
```

Q12. Is this difference statistically significant?

Yes it is

```
t.test(chocolate, fruity)
```

```
##
## Welch Two Sample t-test
##
## data: chocolate and fruity
## t = 6.2582, df = 68.882, p-value = 2.871e-08
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 11.44563 22.15795
## sample estimates:
## mean of x mean of y
## 60.92153 44.11974
```

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
head(candy[order(candy$winpercent),], n=5)
```

```
##               chocolate fruity caramel peanutyalmondy nougat
## Nik L Nip           0      1      0                0      0
## Boston Baked Beans  0      0      0                1      0
## Chiclets            0      1      0                0      0
## Super Bubble        0      1      0                0      0
## Jawbusters          0      1      0                0      0
##
## crispedricewafer hard bar pluribus sugarpercent pricepercent
## Nik L Nip                0  0  0      1      0.197      0.976
## Boston Baked Beans        0  0  0      1      0.313      0.511
## Chiclets                  0  0  0      1      0.046      0.325
## Super Bubble              0  0  0      0      0.162      0.116
```

```
## Jawbusters          0    1    0          1          0.093          0.511
##                               winpercent
## Nik L Nip           22.44534
## Boston Baked Beans  23.41782
## Chiclets            24.52499
## Super Bubble        27.30386
## Jawbusters          28.12744
```

Q14. What are the top 5 all time favorite candy types out of this set?


```
head(candy[order(candy$winpercent, decreasing=TRUE),], n=5)
```

```
##                               chocolate fruity caramel peanutyalmondy nougat
## Reese's Peanut Butter cup      1      0      0              1      0
## Reese's Miniatures             1      0      0              1      0
## Twix                           1      0      1              0      0
## Kit Kat                        1      0      0              0      0
## Snickers                       1      0      1              1      1
##                               crispedricewafer hard bar pluribus sugarpercent
## Reese's Peanut Butter cup              0      0      0      0      0.720
## Reese's Miniatures                    0      0      0      0      0.034
## Twix                                1      0      1      0      0.546
## Kit Kat                             1      0      1      0      0.313
## Snickers                            0      0      1      0      0.546
##                               pricepercent winpercent
## Reese's Peanut Butter cup      0.651  84.18029
## Reese's Miniatures             0.279  81.86626
## Twix                           0.906  81.64291
## Kit Kat                        0.511  76.76860
## Snickers                       0.651  76.67378
```

Q15. Make a first barplot of candy ranking based on winpercent values.

```
library("ggplot2")


ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```

```
#Color vectors (all black to start)
my_cols=rep("black", nrow(candy))
```

```
# Now overwrite the entries with different colors
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```



Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?


Starburst

Taking a look at pricepercent

```
library(ggrepel)

# How about a plot of price vs win
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

```
## Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps
```

Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?


Reeses Miniatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

##	pricepercent	winpercent
## Nik L Nip	0.976	22.44534
## Nestle Smarties	0.976	37.88719
## Ring pop	0.965	35.29076
## Hershey's Kisses	0.918	62.28448
## Hershey's Milk Chocolate	0.918	56.49050

```
# Make a lollipop chart of pricepercent
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy), pricepercent),
                    xend = 0), col="gray40") +
  geom_point()
```




5. Exploring the correlation structure

```
library(corrplot)
```

```
## corrplot 0.90 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity

Q23. Similarly, what two variables are most positively correlated?


Chocolate and winpercent

6. Princial Component Analysis


```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  2.0788 1.1378 1.1092 1.07533 0.9518 0.81923 0.81530
## Proportion of Variance 0.3601 0.1079 0.1025 0.09636 0.0755 0.05593 0.05539
## Cumulative Proportion 0.3601 0.4680 0.5705 0.66688 0.7424 0.79830 0.85369
##          PC8      PC9      PC10     PC11     PC12
## Standard deviation  0.74530 0.67824 0.62349 0.43974 0.39760
## Proportion of Variance 0.04629 0.03833 0.03239 0.01611 0.01317
## Cumulative Proportion 0.89998 0.93832 0.97071 0.98683 1.00000
```

```
plot(pca$x[,1:2])
```




```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])
```

```
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)
```

```
p
```




```
library(ggrepel)

p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
        subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown), fruity (red)",
        caption="Data from 538")
```

```
## Warning: ggrepel: 44 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps
```

Halloween Candy PCA Space


Colored by type: chocolate bar (dark brown), chocolate other (light brown), fruity (red), oth



Data from 538

```
#library(plotly)
#ggplotly(p)
```

```
par(mar=c(8,4,2,2))
barplot(pca$rotation[,1], las=2, ylab="PC1 Contribution")
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard, pluribus. This makes sense because it groups the fruity candy as a different group/cluster characterized by being hard and coming in a bag or box of multiple candies.